

# 1. MNIST

## 1.1 Preliminary results for MNIST

These are results from running the two unmodified tutorials

### MNIST tutorial

Tutorial Title: *TF Layers: Building a CNN MNIST Classifier*

File path in my submission: `other/cnn_mnist_unmodified.py`

**Test accuracy: 0.9699**

Global\_step: 20000

Loss: 0.097417705

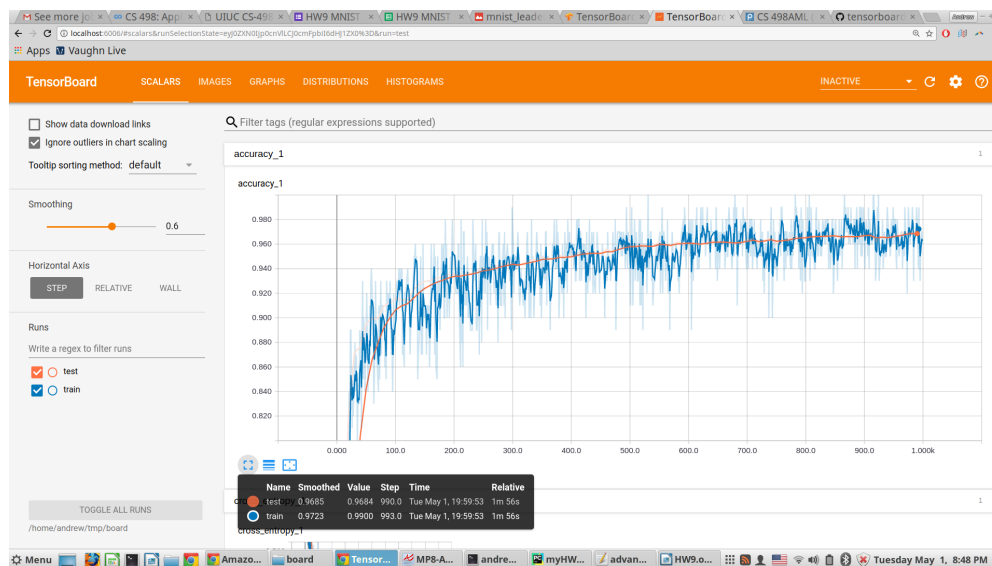
It took me several hours to run this script locally on my machine

### Tensorboard tutorial

Tutorial Title: *TensorBoard: Visualizing Learning*

File path in my submission: `other/mnist_with_summaries_unmodified.py`

**Test accuracy: 0.9683**



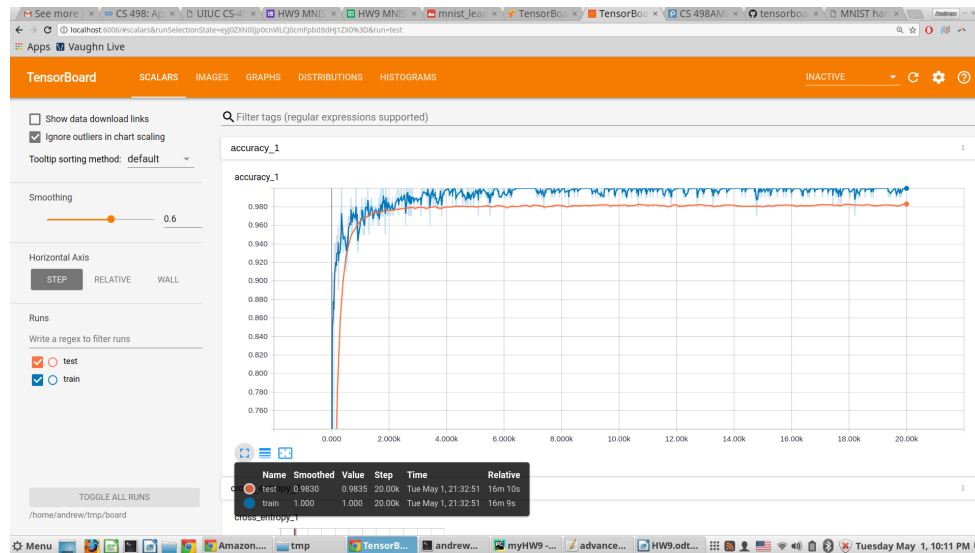
As this was just a tutorial, I ran it for 1000 batches only just to see how it performs.

## Tensorboard tutorial with modifications

Modifications: log the accuracy on tensorboard every 100 batches, for 20,000 batches

File path in my submission: other/mnist\_with\_summaries\_modified.py

Test accuracy: 0.983



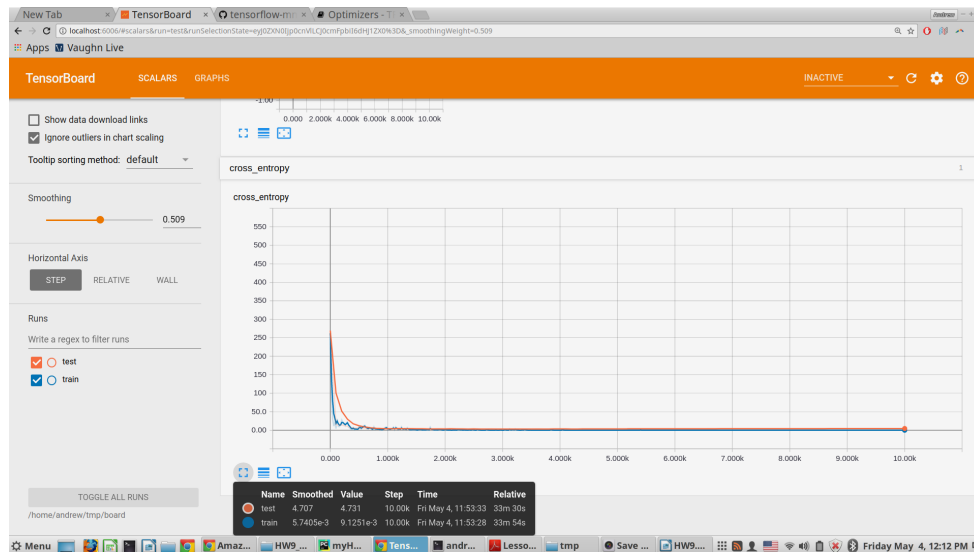
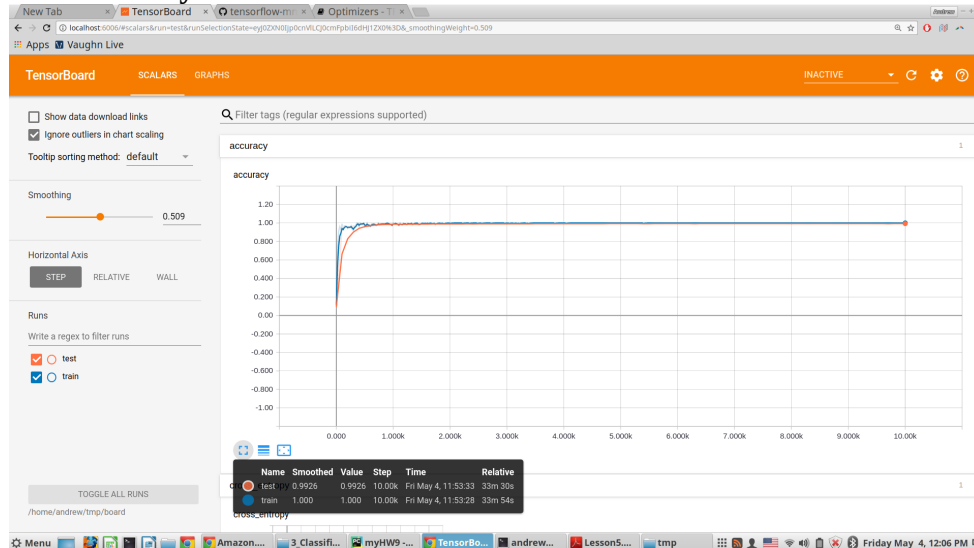
## 1.2 Advanced results for MNIST

**My main solution is based on the Tensorboard tutorial and reference [2]**

Modifications: introduce a convolutional neural network with 3 layers, add Tensorboard functionality (scalars and summaries) in order to visualize results (every 100 batches for 10,000 batches, number of batches was decreased from 20,000 for speed as no significant changes occur after 10,000 in any of the performed runs)

File path in my submission: part1.py

**Test accuracy: 0.9926**



Considering that all these numbers are very close to 100%, the accuracy of 0.9926 is a significant improvement compared to the previous results. It took me several days of test and trial to achieve this locally.

I also tried other approaches similar to reference 2 ([https://github.com/hwalsuklee/tensorflow-mnist-cnn/blob/master/mnist\\_cnn\\_train.py](https://github.com/hwalsuklee/tensorflow-mnist-cnn/blob/master/mnist_cnn_train.py)), but they were very resource-consuming.

## 2. CIFAR-10


### 2.1 Preliminary results for CIFAR-10

As required by the assignment, I made sure I can run the CIFAR-10 tutorial. I did it on a GPU-enabled machine at AWS this time, and it was considerably faster. To expedite the implementation, I decreased the *maximum number of steps* from 1,000,000 to 60,000 in `cifar10_multi_gpu_train.py` and decreased *min\_fraction\_of\_examples\_in\_queue* from 0.4 to 0.2 in `cifar10_input.py`, as it was advised in the CIFAR10 tutorial.

The `cifar10_multi_gpu_train.py` file has code to capture training statistics every 100 batches. In addition, I changed the parameter `eval_interval_secs` in `cifar10_eval.py` from `60*5` to 27 seconds because the average *seconds\_per\_batch* parameter (as displayed in every line output every 10 seconds during training) is 0.27, and this change allowed me to log the eval accuracy for tensorboard for approximately every 100 batches. As I have mentioned, the maximum number of examples was set to 60,000 (instead of the minimum required 2000), so this gave me plenty of data to log on tensorboard. The training script saves a model checkpoint only every 1000 steps which doesn't happen too often, so there is really no need to have the eval script perform evaluation too often. To have more diverse data I changed the point at which model checkpoints are saved from 1000 to 500 in the training script (`cifar10_multi_gpu_train.py`). With this setup, the training and evaluation scripts ran simultaneously (despite the warning in the tutorial) for about **four hours** on AWS's g2.2xlarge instance. Combined, the train and eval event.output files used to construct the below accuracy chart in Tensorboard are about 20 MB in size, therefore I am not attaching them, but they can be provided if needed.

File path in my submission: folder `cifar_baseline`

**Test accuracy: 0.838**



## 1.2 Advanced results for MNIST

When making test runs with the number of batches=2000, I noticed that initial accuracies were pretty high if I changed the batch\_size or the GradientDescentOptimizer to AdamOptimizer. I decided to check the performance on 10,000 batches (this number was selected in the interest of time as the deadline was already approaching) of the following cases:

**Case 1:** AdamOptimizer with the original learning rate = 0.1

batch\_size = 512

One more (third) convolution layer identical to conv2: max\_pool and norm were swapped in conv2 to match conv21

**Accuracy** on 10,000 batches: very poor results

**Case 2:** AdamOptimizer with the initial learning rate = 0.001 and the learning rate decay factor = 0.001 (trick I learned from reference 1, because the original learning rate gave poor performance)

batch\_size = 128

**Accuracy** on 10,000 batches: 0.815

**Case 3:** AdamOptimizer with the initial learning rate = 0.001 and the learning rate decay factor = 0.001,

batch\_size = 512

**Accuracy** on 10,000 batches: 0.828

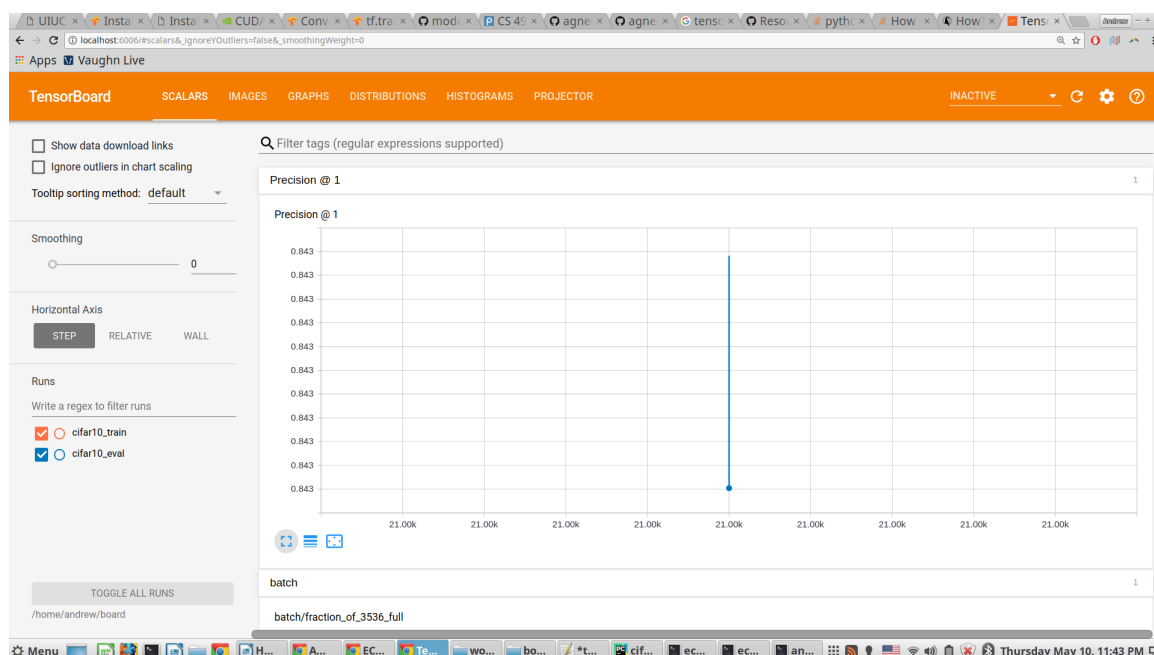
**Case 4:** GradientDescentOptimizer with the original initial learning rate,

batch\_size = 512

**Accuracy** on 10,000 batches: 0.843

As you can see, the best accuracy is in Case 4, but the problem was that the eval script gave me error and did not run together with the train script because of the lack of memory. Therefore, I was able to record only the last final accuracy with the eval script (see below).

File path in my submission: folder cifar\_improved



This is a marginal improvement, but I tried spending days to achieve it. Each run on 10,000 batches took several hours even on the AWS instance. I even rented a more powerful one in the end, but did not help to speed things up significantly.

The next natural step would be to change the CNN architecture. I tried this in Case 1, but was confused by the results which were due to the incorrect learning rate for the AdamOptimizer, as I found out later. To be continued.

## REFERENCES

1. Learning rate schedules. <https://towardsdatascience.com/learning-rate-schedules-and-adaptive-learning-rate-methods-for-deep-learning-2c8f433990d1>
2. Tensorflow CNN MNIST tutorial. <https://github.com/martin-gorner/tensorflow-mnist-tutorial>